

```
In [ ]: import requests
import pandas as pd
url = "https://data.cityofnewyork.us/resource/usep-8jbt.json"

params = {
    "$select": "borough, neighborhood, building_class_category, "
                "land_square_feet, gross_square_feet, sale_price, sale_date",
    "$where": "sale_price > 100000 AND "
                "gross_square_feet IS NOT NULL AND "
                "land_square_feet IS NOT NULL",
    "$limit": 1000
}

resp = requests.get(url, params=params)
df = pd.DataFrame(resp.json())
df
df.to_csv("/Users/lihuixiong/Desktop/DSCI510/FINAL/raw_data.csv", index=False)
```

```
In [4]: numeric_cols = ["sale_price", "land_square_feet", "gross_square_feet"]
for col in numeric_cols:
    df[col] = pd.to_numeric(df[col], errors="coerce")

df["sale_date"] = pd.to_datetime(df["sale_date"], errors="coerce")

df = df.dropna(subset=["sale_price"])
df
```

Out[4]:

		borough	neighborhood	building_class_category	land_square_feet	gross_square_feet	sale_price
0	1	HARLEM-UPPER		11 SPECIAL CONDO BILLING LOTS	0.0	14659	37
1	1	MIDTOWN EAST		11 SPECIAL CONDO BILLING LOTS	0.0	2446	3890
2	1	MIDTOWN EAST		11 SPECIAL CONDO BILLING LOTS	0.0	4321	8500
3	3	DOWNTOWN-FULTON MALL		11 SPECIAL CONDO BILLING LOTS	0.0	129410	11280
4	3	WILLIAMSBURG-SOUTH		11 SPECIAL CONDO BILLING LOTS	0.0	89603	7260
...
995	3	PARK SLOPE		02 TWO FAMILY DWELLINGS	NaN	3800	250
996	3	PARK SLOPE		02 TWO FAMILY DWELLINGS	NaN	3211	211
997	3	WILLIAMSBURG-NORTH		07 RENTALS - WALKUP APARTMENTS	NaN	3218	435
998	3	WILLIAMSBURG-SOUTH		02 TWO FAMILY DWELLINGS	NaN	2730	309
999	5	ARDEN HEIGHTS		01 ONE FAMILY DWELLINGS	NaN	1319	68

1000 rows × 7 columns

In [5]: `df.isna().sum()`

Out[5]:

borough	0
neighborhood	0
building_class_category	0
land_square_feet	993
gross_square_feet	0
sale_price	0
sale_date	0
dtype: int64	

In [6]: `value_counts = df["borough"].value_counts()`

```
for cat, count in value_counts.items():
    print(f'{cat}: {count}')
```

3: 290
5: 258
4: 214
2: 140
1: 98

In [7]: `value_counts = df["building_class_category"].value_counts()`

```
for cat, count in value_counts.items():
    print(f'{cat}: {count}')
```

```
01 ONE FAMILY DWELLINGS: 392
02 TWO FAMILY DWELLINGS: 197
22 STORE BUILDINGS: 57
07 RENTALS - WALKUP APARTMENTS: 53
03 THREE FAMILY DWELLINGS: 52
08 RENTALS - ELEVATOR APARTMENTS: 40
05 TAX CLASS 1 VACANT LAND: 40
29 COMMERCIAL GARAGES: 39
21 OFFICE BUILDINGS: 27
31 COMMERCIAL VACANT LAND: 23
14 RENTALS - 4-10 UNIT: 19
30 WAREHOUSES: 19
37 RELIGIOUS FACILITIES: 8
27 FACTORIES: 6
41 TAX CLASS 4 - OTHER: 5
32 HOSPITAL AND HEALTH FACILITIES: 5
11 SPECIAL CONDO BILLING LOTS: 5
06 TAX CLASS 1 - OTHER: 4
26 OTHER HOTELS: 4
33 EDUCATIONAL FACILITIES: 2
35 INDOOR PUBLIC AND CULTURAL FACILITIES: 2
38 ASYLUMS AND HOMES: 1
```

```
In [9]: def building_class_8cat(cat):
    if cat is None:
        return "Other"
    c = str(cat).upper()

    if c.startswith("01"):
        return "OneFamily"
    elif c.startswith("02"):
        return "TwoFamily"
    elif c.startswith("03"):
        return "ThreeFamily"

    elif c.startswith("07"):
        return "WalkupApt"
    elif c.startswith("08"):
        return "ElevatorApt"

    elif c.startswith("21") or c.startswith("22") or c.startswith("29") or c.startswith("31") or c.startswith("37") or c.startswith("38") or c.startswith("30") or c.startswith("33") or c.startswith("35") or c.startswith("26"):
        return "Commercial"

    elif c.startswith("05") or c.startswith("31") or c.startswith("06"):
        return "VacantLand"

    else:
        return "Other"

df["building_class_8"] = df["building_class_category"].apply(building_class_8cat)
print(df["building_class_8"].value_counts())
```

```

OneFamily      392
TwoFamily     197
Commercial    161
VacantLand     67
WalkupApt      53
ThreeFamily    52
ElevatorApt   40
Other          38
Name: building_class_8, dtype: int64

```

```
In [ ]: df = df.drop(columns=["land_square_feet", "neighborhood"], errors="ignore")
df = df[df["gross_square_feet"] != 0]
df
df.to_csv("/Users/lihuixiong/Desktop/DSCI510/FINAL/processed_data.csv", index=False)
```

	borough	building_class_category	gross_square_feet	sale_price	sale_date	building_class_8
0	1	11 SPECIAL CONDO BILLING LOTS	14659	375000	2025-02-10	Other
1	1	11 SPECIAL CONDO BILLING LOTS	2446	3890000	2025-01-23	Other
2	1	11 SPECIAL CONDO BILLING LOTS	4321	8500000	2025-01-31	Other
3	3	11 SPECIAL CONDO BILLING LOTS	129410	112800000	2024-11-14	Other
4	3	11 SPECIAL CONDO BILLING LOTS	89603	72600000	2025-07-24	Other
...
995	3	02 TWO FAMILY DWELLINGS	3800	2500000	2025-09-04	TwoFamily
996	3	02 TWO FAMILY DWELLINGS	3211	2115000	2025-07-18	TwoFamily
997	3	07 RENTALS - WALKUP APARTMENTS	3218	4350000	2025-03-17	WalkupApt
998	3	02 TWO FAMILY DWELLINGS	2730	3090000	2025-08-07	TwoFamily
999	5	01 ONE FAMILY DWELLINGS	1319	680000	2025-07-18	OneFamily

915 rows × 6 columns

```
In [ ]: import numpy as np
df["log_price"] = np.log(df["sale_price"])
df["log_sqft"] = np.log(df["gross_square_feet"] + 1)

df[["sale_price", "gross_square_feet",
     "log_price", "log_sqft"]].describe()
```

Out[]:


	sale_price	gross_square_feet	log_price	log_sqft
count	9.150000e+02	915.000000	915.000000	915.000000
mean	4.383726e+06	9613.234973	14.302926	8.066457
std	1.235398e+07	23645.223555	1.158157	1.214271
min	1.020000e+05	420.000000	11.532728	6.042633
25%	6.500000e+05	1395.000000	13.384728	7.241364
50%	1.300000e+06	2301.000000	14.077875	7.741534
75%	3.345000e+06	4390.500000	15.022976	8.387426
max	1.640000e+08	238294.000000	18.915377	12.381265

In [17]:

```
import seaborn as sns
import matplotlib.pyplot as plt


numeric_cols = ["sale_price", "gross_square_feet"]
corr = df[numeric_cols].corr()

plt.figure(figsize=(6, 4))
sns.heatmap(corr, annot=True, cmap="Blues", fmt=".2f")
plt.title("Correlation Matrix of Key Numerical Variable")
plt.show()
```



In [19]:

```
plt.figure(figsize=(6, 4))
sns.scatterplot(x=df["gross_square_feet"], y=df["sale_price"])
plt.title("Sale Price vs Gross Square Feet")
plt.show()
```



```
In [ ]: from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import r2_score
from sklearn.ensemble import RandomForestRegressor

# Data Preprocessing
numeric_features = ["log_sqft"]
categorical_features = ["building_class_8", "borough"]

numeric_transformer = Pipeline(steps=[
    ("scaler", StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

# ColumnTransformer
preprocess = ColumnTransformer(
    transformers=[
        ("num", numeric_transformer, numeric_features),
        ("cat", categorical_transformer, categorical_features)
    ]
)

X = df[numeric_features + categorical_features]
y = df["log_price"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```

model_lr = Pipeline(steps=[
    ("preprocess", preprocess),
    ("lr", LinearRegression())
])

model_lr.fit(X_train, y_train)
lr_r2 = model_lr.score(X_test, y_test)
print("Linear Regression R2:", lr_r2)

```

Linear Regression R²: 0.4242983118986853

```

In [ ]: df[ "sale_date" ] = pd.to_datetime(df[ "sale_date" ], errors="coerce")
df[ "year" ] = df[ "sale_date" ].dt.year
df[ "month" ] = df[ "sale_date" ].dt.month
df[ "month_sin" ] = np.sin(2*np.pi*df[ "month" ]/12)
df[ "month_cos" ] = np.cos(2*np.pi*df[ "month" ]/12)

df[ "log_price" ] = np.log(df[ "sale_price" ])
df[ "log_sqft" ] = np.log(df[ "gross_square_feet" ] + 1)

df[ "sqft_borough_interact" ] = df[ "log_sqft" ] * df[ "borough" ].astype("category").cat.
df[ "sqft_class_interact" ] = df[ "log_sqft" ] * df[ "building_class_8" ].astype("category")

numeric_features = [
    "log_sqft",
    "year",
    "month_sin",
    "month_cos",
    "sqft_borough_interact",
    "sqft_class_interact"
]

categorical_features = [ "borough", "building_class_8" ]

numeric_transformer = Pipeline(steps=[
    ("scaler", StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preprocess = ColumnTransformer(
    transformers=[
        ("num", numeric_transformer, numeric_features),
        ("cat", categorical_transformer, categorical_features)
    ]
)

X = df[numeric_features + categorical_features]
y = df[ "log_price" ]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Linear Regression
lr = Pipeline(steps=[
    ("preprocess", preprocess),
    ("lr", LinearRegression())
])
lr.fit(X_train, y_train)
print("LR R2:", lr.score(X_test, y_test))

```

```
LR R2 : 0.4425669849105376
```

```
In [ ]: coef = lr.named_steps["lr"].coef_
intercept = lr.named_steps["lr"].intercept_
print("Intercept:", intercept)
```

```
Intercept: 14.488499333503762
```

```
In [ ]: ohe = lr.named_steps["preprocess"].named_transformers_["cat"]["onehot"]

ohe_feature_names = ohe.get_feature_names_out(categorical_features)

all_features = numeric_features + list(ohe_feature_names)

print("Total features:", len(all_features))
print(all_features)

coef_df = pd.DataFrame({
    "feature": all_features,
    "coefficient": coef
})
```

```
Total features: 18
```

```
[ 'log_sqft', 'year', 'month_sin', 'month_cos', 'sqft_borough_interact', 'sqft_class_interact', 'borough_1', 'borough_2', 'borough_3', 'borough_4', 'borough_5', 'building_class_8_Commercial', 'building_class_8_ElevatorApt', 'building_class_8_OneFamily', 'building_class_8_Other', 'building_class_8_ThreeFamily', 'building_class_8_TwoFamily', 'building_class_8_WalkupApt' ]
```

```
In [33]: coef_df["abs_coef"] = coef_df["coefficient"].abs()
coef_df_sorted = coef_df.sort_values("abs_coef", ascending=False)

print(coef_df_sorted)
```

	feature	coefficient	abs_coef
6	borough_1	1.264093	1.264093
10	borough_5	-1.077126	1.077126
12	building_class_8_ElevatorApt	-0.822455	0.822455
0	log_sqft	0.723789	0.723789
4	sqft_borough_interact	0.398530	0.398530
9	borough_4	-0.378739	0.378739
14	building_class_8_Other	0.307773	0.307773
15	building_class_8_ThreeFamily	0.217013	0.217013
5	sqft_class_interact	-0.176694	0.176694
11	building_class_8_Commercial	0.173645	0.173645
8	borough_3	0.153660	0.153660
16	building_class_8_TwoFamily	0.141824	0.141824
13	building_class_8_OneFamily	-0.089779	0.089779
17	building_class_8_WalkupApt	0.071980	0.071980
2	month_sin	-0.056721	0.056721
7	borough_2	0.038112	0.038112
1	year	0.027334	0.027334
3	month_cos	-0.013495	0.013495

```
In [ ]: import matplotlib.pyplot as plt

topN = 15
coef_df_sorted.head(topN).plot(
    x="feature", y="coefficient", kind="barh", figsize=(8, 6)
)
plt.title("Linear Regression Coefficient Importance")
plt.show()
```

